

# CS-4053 **Recommender System**

Fall 2023

## **Lecture 10:** Generative Adversarial Network (GAN)

**Course Instructor:** Syed Zain Ul Hassan

National University of Computer and Emerging Sciences, Karachi

*Email: [zain.hassan@nu.edu.pk](mailto:zain.hassan@nu.edu.pk)*



# Introduction

- ❑ **Generative Adversarial Networks** (*GANs*) is an unsupervised architecture that use two neural networks, competing against each other (thus the “*adversarial*”) in order to generate new, synthetic instances of data that can pass for real data.
- ❑ GANs can learn to *mimic* any distribution of data.
- ❑ It can be used for image generation, voice generation and video generation tasks.

# Image Generation




[This Person Does Not Exist - Random Face Generator \(this-person-does-not-exist.com\)](http://this-person-does-not-exist.com)




# Introduction

- ❑ A simple GAN architecture consists of two networks, a **generator** and a **discriminator**.



Generator



Discriminator

# Introduction

- ❑ A simple GAN architecture consists of two networks, a **generator** and a **discriminator**.




Generator



Discriminator

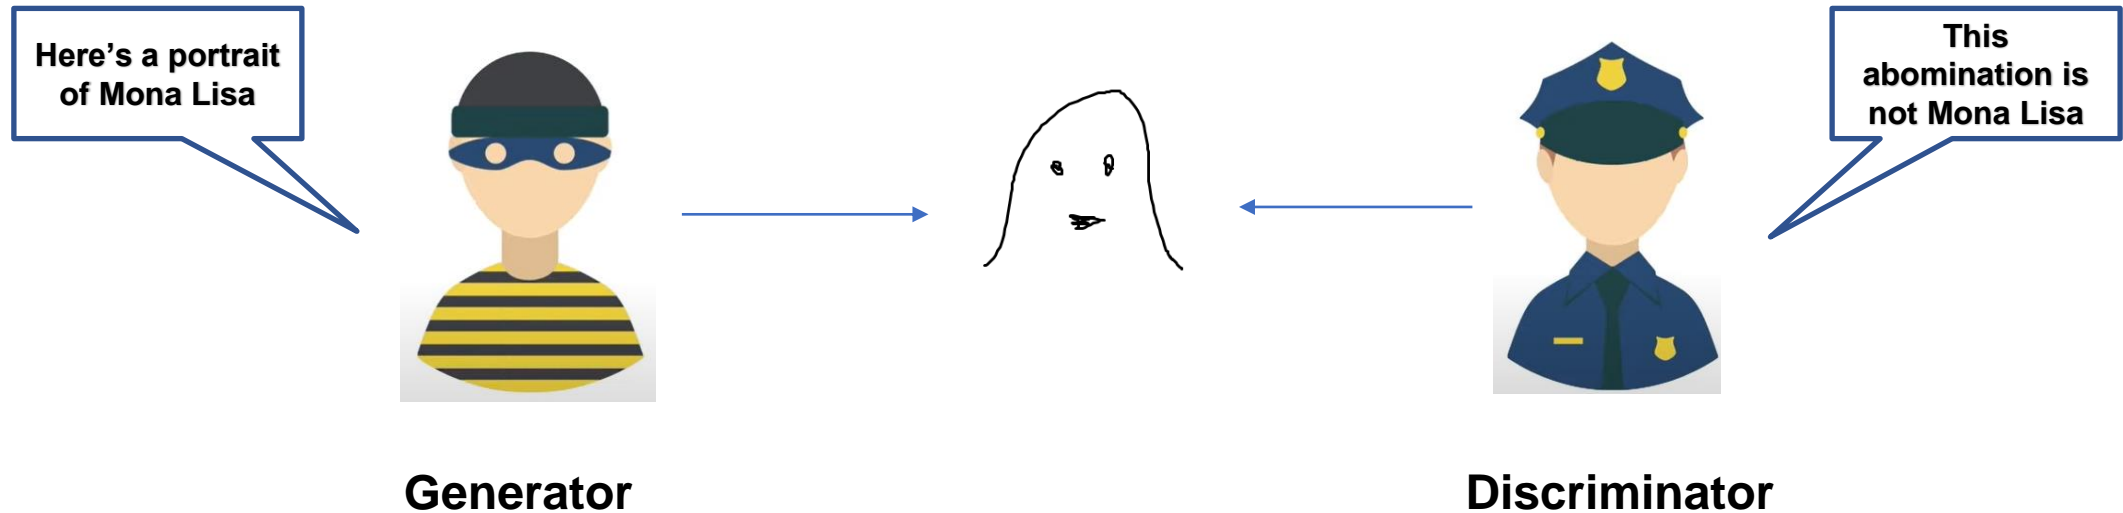
# Adversarial Learning

- ❑ The **generator** starts with random noise and tries to mimic the target distribution (*e.g., image of a digit, image of Mona Lisa*).




# Adversarial Learning

- ❑ The **discriminator** has already seen the training data thus it knows what the target distribution is like thus it acts as a critic.



# Adversarial Learning


- ❑ The **generator** repeatedly tries to fool the discriminator into believing that the generated distribution (image) is a real one.






# Adversarial Learning

- ❑ Based on the feedback (*loss*) of the **discriminator** the quality of generated output is gradually improved.




# Adversarial Learning

- ❑ This adversarial learning is a two-player game between the **generator** and the **discriminator**.



# Adversarial Learning

- This adversarial learning is a two-player game between the **generator** and the **discriminator**.



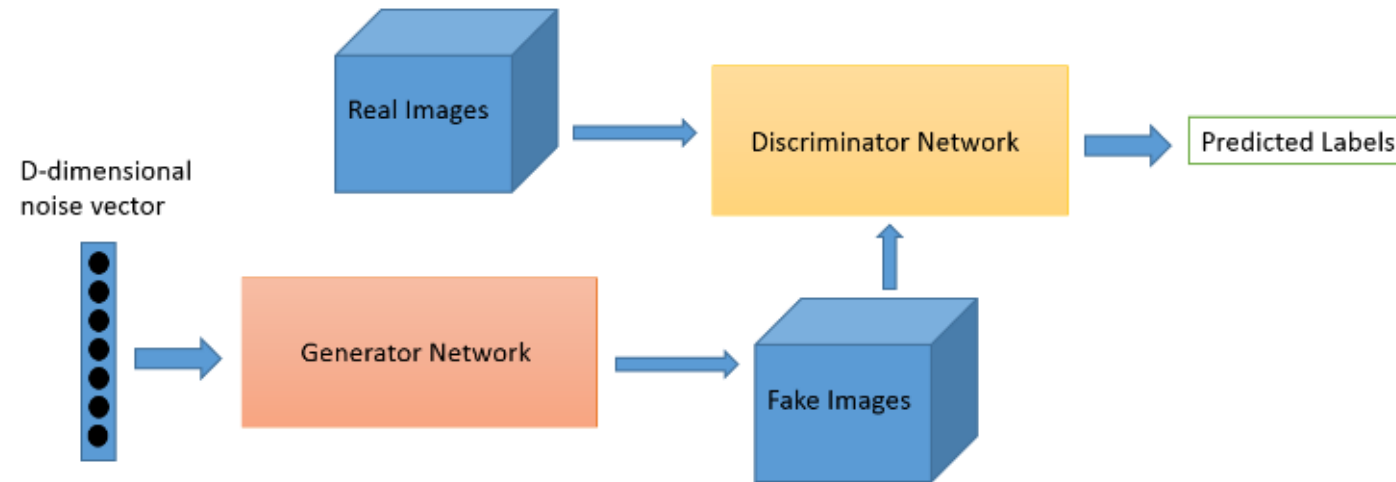
# Steps taken by GAN:

- *The generator takes in random numbers and returns an image.*
- *This generated image is fed into the discriminator alongside a stream of images taken from the actual, ground-truth dataset.*
- *The discriminator takes in both real and fake images and returns probabilities, a number between 0 and 1, with 1 representing a prediction of authenticity and 0 representing fake.*

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

# Steps taken by GAN:

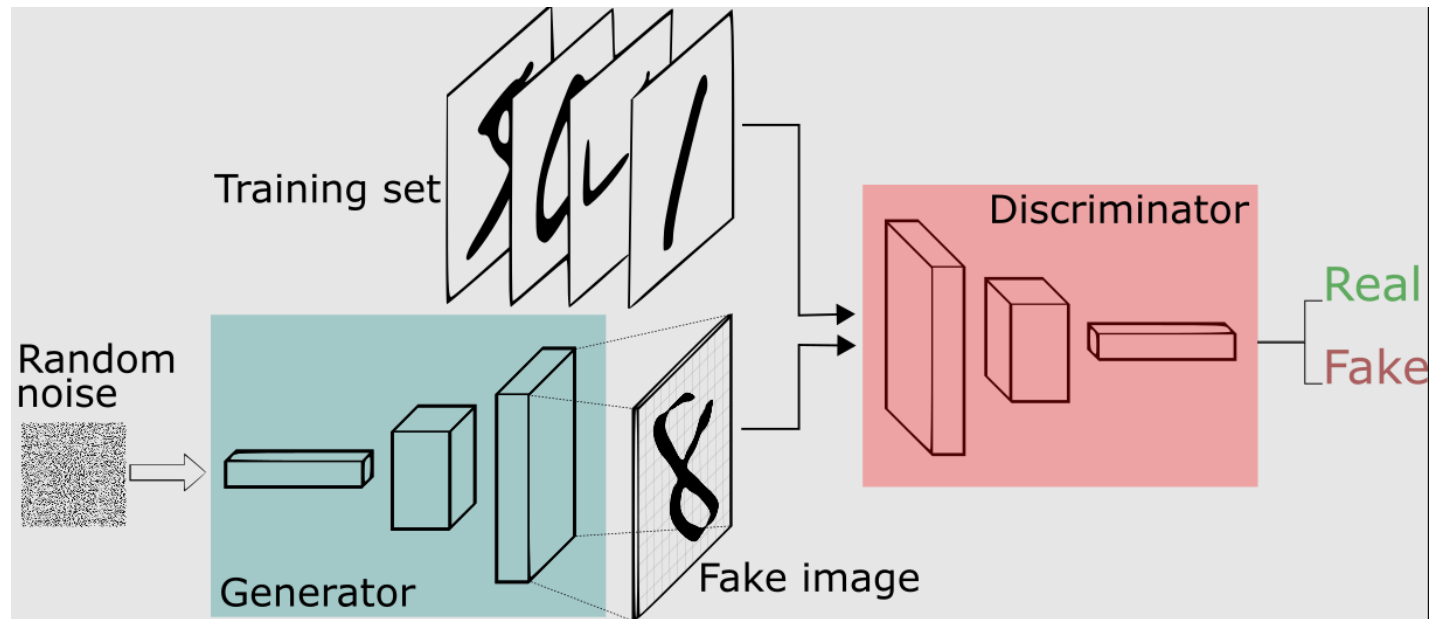


*Image Credits: O'Reilly*

# Example: **Generating MNIST digits**

- ❑ For MNIST, the discriminator network is a standard convolutional network that can categorize the images fed to it, a binomial classifier labeling images as real or fake.
- ❑ The generator is an inverse convolutional network that takes a vector of random noise and upsamples it to an image

# Example: **Generating MNIST digits**



# How to train a GAN

- ☐ Build your generator and discriminator models
- ☐ Sample data from true distribution and train discriminator on this data
  - ☐ Discriminator works as a binary classifier
- ☐ Create input data from latent space (noise) and feed it to the generator
- ☐ The generator output serves as input to the discriminator
- ☐ The discriminator uses Sigmoid activation to classify the generated output as real or fake
- ☐ The discriminator loss is fed-back to the generator
- ☐ The generator is trained until it reaches convergence



# Challenges

- ❑ GANs are really hard to train!
- ❑ Hard to meet convergence criteria.
- ❑ Vanilla GANs suffer from vanishing gradient problem.
- ❑ Overwhelming adversary i.e. one network should not “*overpower*” another
- ❑ Cannot count objects (*Stable Diffusion can help*)
- ❑ Sensitive to choice of hyperparameters
- ❑ ...and many other challenges

# Guidelines for training

- ☐ Use ***tanh*** activation for **generator** and ***sigmoid*** for **discriminator**
- ☐ Freeze **generator** for every other training iteration of **discriminator**
- ☐ Use ***LeakyReLU*** in dense layers
- ☐ Use Functional API (for Keras)
- ☐ Use dropout

# Types of GANs

- ☐ Vanilla GANs (original model)
- ☐ CGAN
- ☐ Wasserstein GAN
- ☐ StyleGAN
- ☐ InfoGAN
- ☐ DCGAN
- ☐ DualGAN
- ☐ KDGAN
- ☐ ... *many other variants*

# Which architecture to use for RecSys?

- ☐ Vanilla GANs (original model)
- ☐ **CGAN**
- ☐ Wasserstein GAN
- ☐ **StyleGAN**
- ☐ InfoGAN
- ☐ DCGAN
- ☐ DualGAN
- ☐ KDGAN
- ☐ ... *many other variants*

# Which architecture to use for RecSys?

- ☐ Vanilla GANs (original model)
- ☐ **CGAN**
- ☐ Wasserstein GAN
- ☐ **StyleGAN**
- ☐ InfoGAN
- ☐ DCGAN
- ☐ DualGAN
- ☐ KDGAN
- ☐ ... *many other variants*